

# Fundamentals of Database Systems

## Concurrency Control Protocols

Arnab Bhattacharya  
Dept. of Computer Science and Engineering,  
Indian Institute of Technology, Kanpur

NPTEL  
[https://onlinecourses.nptel.ac.in/noc15\\_cs14/](https://onlinecourses.nptel.ac.in/noc15_cs14/)

July-September, 2015

- Locks are used to control access to a data item
- Lock requests are made to **concurrency control manager**
- Concurrency control manager decides whether and when to grant locks
- Locking and unlocking must be **atomic**
- A data item can be locked in two modes
  - ① **Exclusive (X)** mode: Data item can be both written and read
  - ② **Shared (S)** mode: Data item can only be read
- A lock can be granted based on the compatibility matrix
- **Lock compatibility matrix** or **conflict matrix**

	S	X
S	yes	no
X	no	no

- If a lock cannot be granted, it must wait

# Locking protocol

- A schedule must specify all locking and unlocking operations and their modes
  - $lx(a)$  requests an exclusive lock on data item  $a$ ;  $ux(a)$  releases it
  - $ls(a)$  requests a shared lock on data item  $a$ ;  $us(a)$  releases it
- Example:  $lx_1(a); r_1(a); w_1(a); ls_2(b); r_2(b); ux_1(a); us_2(b)$
- Consider  $lx_1(a); r_1(a); w_1(a); ls_2(b); r_2(b); lx_2(a); lx_1(b)$ 
  - **Deadlock**
- **Starvation** may also happen
- A **locking protocol** specifies the rules of how a transaction can acquire and release locks

# Two-phase locking protocol (2PL)

- Two phases
- Phase 1: **Growing (locking) phase**
  - Transaction may obtain locks
  - Transaction may not release locks
- Phase 2: **Shrinking (unlocking) phase**
  - Transaction may release locks
  - Transaction may not obtain locks
- 2PL schedules are conflict serializable
  - Serialized in the order of **lock points**
  - Lock point is the time when *all* locks are obtained
- May suffer from deadlock
  - $lx_1(a); r_1(a); w_1(a); ls_2(b); r_2(b); ls_2(a); lx_1(b)$
- May suffer from cascading rollbacks
  - $lx_1(a); r_1(a); w_1(a); ux_1(a); lx_2(a); r_2(a); w_2(a); ux_2(a); ls_3(a); r_3(a); a_1$

# Variants of 2PL

- Basic 2PL

- Basic protocol

- Strict 2PL

- A transaction must hold all its *exclusive* locks till it commits or aborts
- Avoids cascading rollbacks
- Produces strict schedules
- May deadlock

- Rigorous 2PL

- A transaction must hold *all* its locks till it commits or aborts
- Transactions can be serialized in the order of their commits

- Conservative (static) 2PL

- All locks are acquired atomically before a transaction begins
- Each transaction declares its **read set** and **write set**
- Deadlock-free

# Timestamps

- Each transaction is assigned a **timestamp** when it starts
  - Transaction  $T_i$  starting earlier has a lower timestamp than  $T_j$  starting later
- For each data item  $x$ , two timestamps are maintained
- **write-timestamp(x)** is the largest timestamp of any transaction that executed write successfully
- **read-timestamp(x)** is the largest timestamp of any transaction that executed read successfully
- Protocols using timestamps *cannot* deadlock

# Timestamp ordering (TO) protocol

- Ensures that conflicting operations are executed in timestamp order
- When a transaction  $T$  requests  $\text{read}(x)$ 
  - If  $\text{ts}(T) < \text{wts}(x)$ , then  $T$  is attempting to read a value that has been overwritten by a later transaction
    - $\text{read}(x)$  request is rejected
    - $T$  or transaction that produced  $\text{wts}(x)$  is rolled back
  - If  $\text{ts}(T) > \text{wts}(x)$ , no conflict
    - $\text{read}(x)$  request is executed
    - $\text{rts}(x)$  is updated to  $\max\{\text{rts}(x), \text{ts}(T)\}$
- When a transaction  $T$  requests  $\text{write}(x)$ 
  - If  $\text{ts}(T) < \text{rts}(x)$ , then value of  $x$  that is being written should have been read earlier
    - $\text{write}(x)$  request is rejected
    - $T$  or transaction that produced  $\text{rts}(x)$  is rolled back
  - If  $\text{ts}(T) < \text{wts}(x)$ , then  $T$  is attempting to write an obsolete value that has been overwritten by a later transaction
    - $\text{write}(x)$  request is rejected
    - $T$  or transaction producing  $\text{wts}(x)$  is rolled back
  - If  $\text{ts}(T) > \text{rts}(x)$  and  $\text{ts}(T) > \text{wts}(x)$ , no conflict
    - $\text{write}(x)$  request is executed
    - $\text{wts}(x)$  is updated to  $\text{ts}(T)$

- Guarantees conflict serializability
- Conflicting operations are executed in timestamp order
  - If an operation appears out of order, it is rejected
- No deadlock since all edges in the precedence graph are from transactions with smaller timestamp to those with larger timestamp
- May cause starvation
- Is not cascadeless
- Is not recoverable



# Modifications

- To make it recoverable
  - Use **commit dependency**
  - If  $T_i$  reads from  $T_j$  and  $T_j$  has not committed, then  $T_i$  has a commit dependency on  $T_j$
  - Ensure that  $T_i$  does not commit before  $T_j$  commits
- To make it recoverable and cascadeless
  - All writes are performed *atomically* in the end
  - A transaction, if aborts, is re-started with a new timestamp
- To make it recoverable and cascadeless
  - Lock data that is begin written
  - Wait for it to be committed before allowing read
- **Strict timestamp ordering**: to make it strict
  - Wait for data to be committed before reading or writing

# Thomas' write rule

- Obsolete writes may be ignored
- When T attempts to write x, if  $ts(T) < wts(x)$ , then T is trying to write an obsolete value of x
- **Thomas' write rule**: Rather than aborting T, ignore the write operation
  - Write is obsolete anyway
- Improves concurrency and recoverability
- Allows some view-serializable schedules that are not conflict-serializable
  - $r_1(a)w_2(a)w_1(a)w_3(a)$

# Validation (certification)-based protocol

- Three phases of a transaction T
- **Read and execution phase**: T writes only to local temporary variables
- **Validation phase**: T performs **validation test** to determine if local variables can be written without violating serializability
- **Write phase**: If T is validated, it updates the database; otherwise it is rolled back (actually nothing needs to be done)
- Also known as **optimistic concurrency control** since transaction executes fully in the hope that all is well
- Three timestamps for each transaction
  - $start(T)$ : start of execution phase
  - $validation(T)$ : start of validation phase
  - $finish(T)$ : end of write phase
- Timestamp of T is set to validation timestamp:  $ts(T) = validation(T)$
- Serialized using this timestamp
  - Increases concurrency
- Cascadeless
- Starvation
- No deadlock


# Validation test

- For a transaction  $T_i$ , check two conditions for all transactions  $T_j$  with  $ts(T_j) < ts(T_i)$ 
  - $finish(T_j) < start(T_i)$
  - $finish(T_j) < validation(T_i)$  and the read-set of  $T_i$  is disjoint from the write-set of  $T_j$
- If either of these conditions is true, validation succeeds; otherwise, it fails
- Justification
  - First condition ensures serial schedules
  - Writes of  $T_i$  cannot affect reads of  $T_j$
  - Writes of  $T_j$  do not affect reads of  $T_i$  as they are disjoint

# Mutiple granularity

- Hierarchy of data items
  - DB, Relation, Tuple, Attribute
- Locking can be done at different levels
- Locking a node *explicitly* locks all its descendants *implicitly*
  - Explicit locks
  - Implicit locks
- Granularity of locking
  - *Fine granularity*: lower in tree, high concurrency, high locking overhead
  - *Coarse granularity*: higher in tree, low concurrency, low locking overhead

# Problems of simple locking



- Assume that  $T_1$  has locked tuple  $f_2$
- $T_2$  wants to lock  $a_{24}$ 
  - It cannot since  $a_{24}$  is implicitly locked
  - Find out by traversing path from  $a_{24}$  to  $d$
- $T_3$  wants to lock  $r_1$ 
  - It cannot since that would lock  $f_2$  implicitly
  - Find out by searching entire subtree under  $r_1$
- Thus, for efficiency, **intention lock modes** are used
  - Ancestors of an explicitly locked node are in intention mode

# Intention lock modes

- In addition to **shared (S)** and **exclusive (X)** locks, three additional locks
- **Intention-shared (IS)**: at least one descendant has a S lock
- **Intention-exclusive (IX)**: at least one descendant has a X lock
- **Shared and intention-exclusive (SIX)**: node is locked in S mode *and* at least one descendant has X lock

# Rules of locking

- A transaction may obtain only one lock on an entity at a time
- If two locks are needed, the more restrictive one will be acquired
- Every lock must be given notice of all lower-level locks
- Locks are acquired in root-to-leaf order
- Locks are released in leaf-to-root order
- Compatibility matrix

	IS	IX	S	SIX	X
IS	yes	yes	yes	yes	no
IX	yes	yes	no	no	no
S	yes	no	yes	no	no
SIX	yes	no	no	no	no
X	no	no	no	no	no



# Multiple granularity locking scheme

- Transaction T wants to lock a node x:
  - Lock compatibility matrix is observed
  - In S or IS mode: only if parent of x is locked by T in IX or IS mode
  - In X, SIX or IX mode: only if parent of x is locked by T in IX or SIX mode
  - Maintains 2PL, i.e., has not unlocked anything
- Transaction T wants to unlock a node x:
  - No child of x is currently locked by T
- Ensures conflict serializability

- Suppose  $T_1$  wants to read  $r_1$  but only modify  $a_{24}$
- Locking  $r_1$  in IX mode will allow other transactions to lock  $r_1$  in IX mode
  - Unsafe as  $T_1$  is reading  $r_1$
- Locking  $r_1$  in S mode will allow other transactions to lock  $r_1$  in S mode and read everything
  - Unsafe as  $T_1$  is modifying  $a_{24}$
- SIX lock compromises and is safer

# Example

- $T_1$  wants to read  $a_{12}$ 
  - Locks  $d, r_1, f_1$  in IS mode and  $a_{12}$  in S mode
- $T_2$  wants to write  $a_{14}$ 
  - Locks  $d, r_1, f_1$  in IX mode and  $a_{12}$  in X mode
- $T_3$  wants to read  $f_1$ 
  - Locks  $d, r_1$  in IS mode and  $f_1$  in S mode
- $T_4$  wants to read  $d$ 
  - Locks  $d$  in S mode
- $T_1$  and  $T_2$  can execute concurrently
- $T_1, T_3$  and  $T_4$  can execute concurrently
- $T_2$  and  $T_3$  cannot execute concurrently
- $T_2$  and  $T_4$  cannot execute concurrently

# Deadlock prevention

- Deadlock prevention schemes *never* allow a system to enter deadlock
- Two schemes that use timestamps
- **Wait-die**: Non-preemptive
  - Older transactions wait for younger ones to release data item (**wait**)
  - Younger ones do not wait, they roll back (**die**)
  - Young transactions may die many times
- **Wound-wait**: Preemptive
  - Older transactions kill younger ones and force them to release data item (**wound**)
  - Younger ones wait (**wait**)
- Transactions are re-started with the *same* timestamp
- No starvation
- Wound-wait has fewer rollbacks than wait-die
  - Less likely for old transactions to not finish and want a lock from a young transaction

# Deadlock recovery

- Deadlocks can be detected by a **wait-for** graph
  - $T_i$  waits for  $T_j$ ,  $T_j$  waits for  $T_k$ , etc.
- If deadlock is detected by finding a cycle, a transaction must be chosen for roll back, i.e., made a **victim**
- Which transaction?
  - One with lowest cost
  - One with least progress
  - One inducing least number of cascading rollbacks
- How far to rollback?
  - **Total rollback**: Completely abort and re-start
  - **Partial rollback**: Rollback to only as far as necessary to break deadlock
- Starvation happens when the same transaction is repeatedly chosen as the victim
  - Factor number of rollbacks when choosing victim

# Insert and delete

- **insert(x)**: inserts the data item x
- **delete(x)**: deletes the data item x
- Logical errors
  - read(x), write(x) before insert(x)
  - read(x), write(x) after delete(x)
  - delete(x) after delete(x)
  - insert(x) after insert(x)
- Conflicts
  - Similar to write(x)

# Phantom phenomenon

- Transaction T1 reads an entire relation to compute an aggregate
- Transaction T2 inserts tuples to the relation
- Transactions T1 and T2 *logically conflict*
- However, they do not conflict on any tuple
- They conflict on a **phantom tuple**
  - This is called **phantom phenomenon**
- Exclusive lock on relation solves this
- Multi-granularity locking
- If index structure is used, **index locking** protocol improves concurrency by locking index nodes
  - Avoids phantom phenomenon since every transaction needs to lock all accessed nodes