

\* Semaphores: It gives a 'n' process soln.  
 "A Semaphore  $s$  is an integer variable that, open from initialization, is accessed only through 2 standard atomic operations:  $\text{wait}()$  &  $\text{signal}()$ .

Initialization :  $\text{int } s=1;$

wait(): is an atomic operation which tests the condition & decrements the value of  $s$  by 1.

signal(): is an atomic operation which increments the value of  $s$  by 1.

Semaphores will solve : 1) CS problem

Applications 2) Order of execution among processes  
 3) Resource mgmt.

$\text{int } s=1;$

Consider  $n$  processes  $P_1, P_2, P_3, \dots, P_n$

do	$\text{wait}(s)$	$\text{signal}(s)$
{	{	{
$\text{wait}(s);$	$\text{while}(s \leq 0);$	$s = s + 1;$
[CS]	$s = s - 1;$	
$\text{signal}(s);$	}	
[RS]		
}		
while(true)		

✓ Let's consider that initially process  $P_1$  is executing, it will execute  $\text{wait}()$  fun. & since the cond.  $\text{while}(s \geq 0)$  is false it will come out of loop & decrements the value of  $s$  (i.e.,  $s=0$  now). &  $P_1$  enters in its CS. Now if  $P_2$  wants to enter into CS, it has to execute  $\text{wait}()$  operation but the condition  $\text{while}(s \geq 0)$  becomes true & hence  $P_2$  will have to wait until the condition becomes

false. i.e., it is unable to enter into CS. similarly until PI comes out of CS & increments the S value by 1, no other process can enter into its CS.



① Mutual Exclusion is guaranteed

② Progress = ?

Here, only those processes that are interested in entering CS only will execute/check the wait() operation & they may enter. We are not restricting if the processes are not interested.



Progress is guaranteed

③ Bounded Wait ?

Whenever a process comes out of CS, any other process/some process can decrement the value of S and can enter the CS. So hence we can't determine a logical bound time for any process because the entry is basically determined by the OS, randomly (i.e., whichever the process is having control over the CPU will enter the CS).



Bounded Wait is not guaranteed.

✓ MF & Progress are the mandatory criterias which are satisfied by semaphores. But it will not satisfy bounded wait, which is not essential. Hence we can call this as a soln to CS.

There are 2 types of Semaphores:

- 1) Binary Semaphores: can range only between 0 & 1
- 2) Counting Semaphores: can range over unrestricted domain

\* On some systems, binary semaphores are known as mutex locks, as they are locks that provide mutual exclusion.

```
do
{ wait(mutex);
  [C9]
  Signal(mutex);
  [RS]
} while(true);
```

\* On binary semaphores mutex is initialized to 1.

### \* Classic Problems of Synchronization:

#### 1) Producer-Consumer Problem / Bounded-Buffer Problem:

We assume that the pool consists of  $n$  buffers, each capable of holding 1 item. The mutex semaphore provides mutual exclusion for accesses to the buffer pool & is initialized to the value 1. The empty & full semaphores count the no. of empty & full buffers. The semaphore empty is initialized to  $n$  & full is initialized to 0.

mutex = 1;

empty =  $n$ ;

full = 0;

#### Producer Process

```
do
{
  ...
  // produce an item
  ...
  wait(empty);
  wait(mutex);
  ...
  // add item to buffer
```

#### Consumer Process

```
do
{
  ...
  wait(full);
  wait(mutex);
  ...
  // Remove an item from buf
  ...
  Signal(mutex);
```

```

    signal(mutex);
}
signal(full);
fwhile(true);
    signal(empty);
    ...
    // consume the item
    ...
}while(true);

```

Producer: Produces the items & adds them to the buffer.

Producer will decrement empty value & checks mutex condition, & produces the item & increments the value of full & then releases mutex (decrements by 1).

✗ Producer has to check the overflow condition: He can produce the items only if atleast 1 buffer is empty. If empty = 0  $\Rightarrow$  condition becomes true & waits until empty is incremented by the consumer.

Consumer: Consumes the items & deletes them from the buffer. consumer will decrement the value of full & checks mutex condition & enters cs to consume the item. After consuming, consumer will increment mutex as well as empty value.

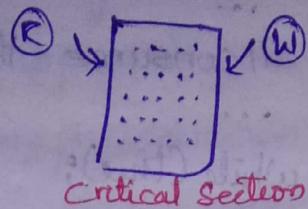
consumer has to check the underflow condition. He can consume items only if atleast 1 buffer/item is available. If full = 0  $\Rightarrow$  condition becomes true & hence consumer waits until producer increments the value of full to atleast 1.

d) Reader-Writer Problems: Lets consider a database or file & 2 types of processes  $\rightarrow$  Reader process + writer process.

Reader process: Read the file or database

Writer process: can do both - Reading & Writing.

Here the file or db is the critical section.



Critical Section

- ✓ For Readers it's actually not a CS because more than one reader can read at the same time & hence more than 1 reader can enter into CS. But it is CS for readers, only when writer wants to enter into CS. Even if 1 reader is present in the CS, no writer shld be allowed to enter the CS.
- ✓ For writers its a CS in both the ways. When a writer is present in the CS, we shld not allow either a reader or another writer to enter into CS.
- ✓ Let's solve it using semaphores: Consider 2 Semaphores i.e., wrt + mutex.

For Writer

```
while(1)
{
    Wait (wrt);
    // Write operation
    signal (wrt);
}
```

Initialization:

```
wrt = 1
mutex = 1
readcount = 0
```

For Reader

```
while(1)
{
    wait (mutex);
    readcount++;
    if (readcount == 1)
        wait (wrt);
    signal (mutex);
}

// Read operation
wait (mutex);
readcount--;
if (readcount == 0)
    signal (wrt);
signal (mutex);
```

\* Writer: We have used semaphore wrt = 1, whenever writer wants to perform write operation, wait(wrt) is performed & value becomes 0, until writer increments the value

of wrt/complete the write operation, no other Reader/Writer process is allowed to enter cs.

✓ Very easy.

For Reader: If reader is present in the cs, & any other process wants to enter into cs, we shld check whether its a reader process or a writer process. If it is a reader process, we can allow it to enter cs but if it is a writer process, we shld not allow.

✓ We shld also note that whenever the 1<sup>st</sup> reader enters into cs, it is the responsibility of the ~~first~~ reader to stop writer processes to enter into cs. In the similar fashion, the last reader process, which is exiting from the cs shld allow the writer to enter into cs.

readcount = 0 (initially)

Mutex is used to get synchronization b/w reader processes i.e., only 1 reader process shld increment the readcount or decrement the readcount at any point of time.

3) The Dining - Philosophers Problem: Consider 5 philosophers, who spend their lives thinking & eating. The philosophers share a circular table surrounded by 5 chairs, each belonging to 1 philosopher. In the centre of the table is a bowl of rice & the table is laid with 5 single chopsticks. When a philosopher thinks, he does not interact with his colleagues. From time to time, a philosopher gets hungry & tries to pick up the 2 chopsticks that are closest to him (the chopsticks that are b/w his left & right neighbours.).

✓ A philosopher may pick up only 1 chopstick at a time. Obviously, he can't pickup a chopstick that is already in hand of a neighbour.

✓ When a hungry philosopher have both chopsticks at the same time, he eats without releasing her chopsticks. When she is finished eating, she puts down both of his chopsticks & starts thinking again.

✓ One simple sol<sup>n</sup> to represent each chopstick with a semaphore. A philosopher tries to grab a chopstick by executing a wait() operation on that semaphore; she releases his chopsticks by executing the signal() operation on the appropriate semaphore. Thus the shared data are:

Semaphore chopstick[5];

where all the elements of chopsticks are initialized to 1.

The structure of philosopher "i" is shown below:

Sol<sup>n</sup> 1:

do

{

wait(chopstick[i]);

wait(chopstick[(i+1)%5]);

...

// eat

...

signal(chopstick[i]);

signal(chopstick[(i+1)%5]);

...

// think

...

}while(True);

✓ The above Sol<sup>n</sup> guarantees that no 2 philosophers are eating simultaneously but it is rejected because it could create a deadlock. Suppose that all 5 philosophers are hungry, each grabs his left chopstick. All the 5 elements of chopsticks will now be equal to zero.

And deadlock occurs.

Sol<sup>n</sup> 2: Allow atmost 4 philosophers to pick up left chopstick first & other philosopher picks up right chopstick first.

<del>Lefty</del>	<del>Righty</del>
<pre>     wait(chopstick[i]);     wait(chopstick[(i+1)%5]);     //eat     signal(chopstick[i]);     signal(chopstick[(i+1)%5]);     //think   </pre>	<pre> do {   wait(chopstick[(i+1)%5]);   wait(chopstick[i]);   //eat   signal(chopstick[(i+1)%5]);   signal(chopstick[i]);   //think } while(true);   </pre>

This will solve the deadlocks problem, as atleast 1 philosopher will get 2 chopsticks.

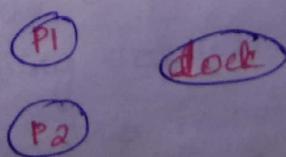
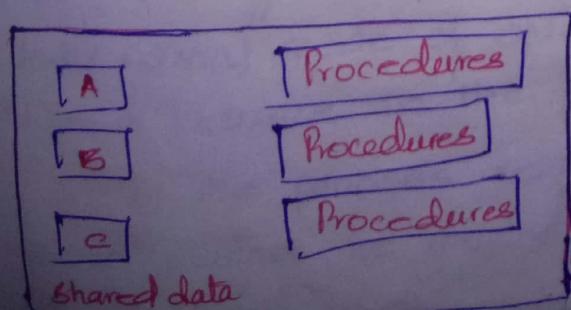
\*Monitors: Although Semaphores provides a convenient & effective mechanism for process synchronization, using them incorrectly can result in timing errors that are difficult to detect, since these errors can happen only if some particular execution sequences take place & these sequence do not always occur.

✓ A monitor is an abstract data type which presents

a set of programmer defined operations that are provided mutual exclusion within the monitor.

✓ It encapsulates private data with public methods to operate on that data.

Monitor



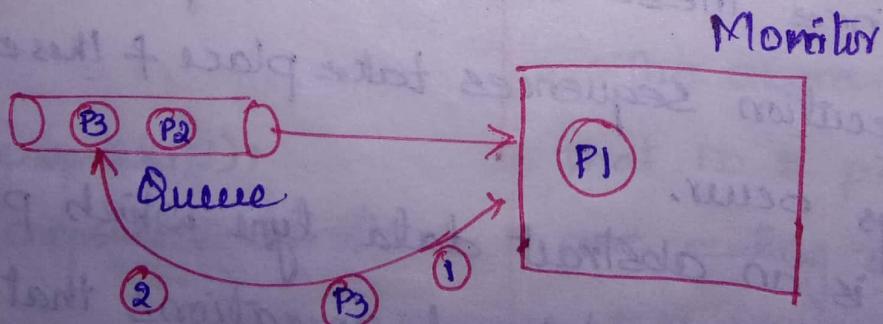
✓ Monitors contains shared data & also procedures which needs to access data. Processes will not directly access data. Processes have to call procedures & procedures in turn allow access to data. And only one process can access data at any point of time.

✓ Monitors are associated with locks & processes which have the lock can only access data.

✗ Monitor is a module that encapsulates:

- 1) Shared data structures
- 2) Procedures that operate on the shared data
- 3) Synchronization b/w concurrent procedure invocation.

✗ If process P<sub>1</sub> wants to access shared data it will contact the procedure & gets a lock & enters CS now other processes will not be allowed to enter the monitor.



## 7. Deadlocks

✓ In a multiprogramming environment, several processes may compete for a finite no. of resources. A process requests a resource; if all the resources are not available at that time, the process enters a waiting state. Sometimes a waiting process is never again able to change state, because the resources it has requested are held by other waiting processes. This situation is called as a deadlock.

System Model: A System consists of finite no. of resources to be distributed among a no. of competing processes. The resources are partitioned into several types, each consisting of some no. of identical instances. Memory Space, CPU cycles, files, I/O devices (printers, DVD driver) are examples of resource types. If a system has 2 CPU's, then the resource type CPU has 2 instances. If a process requests an instance of a resource type, the allocation of any instances of the type will satisfy the request.

✓ A process must request a resource before using it and must release the resource after using it. The no. of resources required by a process should not exceed the total no. of resources available in the system.

✓ Under normal mode of operation, a process may utilize a resource in the following sequence:

1) Request: The process requests the resource. If the resource cannot be granted immediately (for e.g. it is being used by another process), then the requesting

process must wait until it can acquire the resource.

2) Use: The process can operate on the resource.

3) Release: The process releases the resource.

✓ A programmer who is developing multithreaded applications must pay particular attention to this problem. Multithreaded progs are good candidates for deadlock because multiple threads can compete for shared resources.

Deadlock characterization: Features that characterize deadlocks.

Necessary conditions: A deadlock situation can arise if the following 4 conditions hold simultaneously in the system.

1. Mutual Exclusion: Atleast 1 resource must be held in a nonshareable mode; that is only 1 process at a time can use the resource. If another process requests the resource, the requesting process must be delayed until the resource has been released.

2. Hold and Wait: A process must be holding atleast 1 resource & waiting to acquire additional resources that are currently being held by other processes.

3. No Preemption: Resources cannot be preempted; that is, a resource can be released only voluntarily by the process holding it, after that process has completed its task.

4. Circular Wait: A set  $\{P_0, P_1, \dots, P_n\}$  of waiting processes must exist such that  $P_0$  is waiting for a resource held by  $P_1$ ,  $P_1$  is waiting for a resource held by  $P_2, \dots, P_{n-1}$  is

Waiting for a resource held by  $P_n$  &  $P_n$  is waiting for a resource held by  $P_0$ .

We emphasize that all 4 conditions must hold for a deadlock to occur. These conditions are not completely independent.

Resource-Allocation Graph: Deadlocks can be described more precisely in terms of a directed graph called a system Resource-Allocation graph.

This graph consists of a set of vertices  $V$  and set of edges  $E$ . The set of vertices  $V$  is partitioned into 2 different types of nodes:  $P = \{P_1, P_2, P_3, \dots, P_n\}$ , the set consisting of all the active processes in the system, &  $R = \{R_1, R_2, \dots, R_m\}$ , the set consisting of all resource types in the system.

Notations:  
①  $P_i \rightarrow R_j$ : Process  $P_i$  has requested an instance of resource type  $R_j$  & is currently waiting for the resource (Request Edge).

②  $R_j \rightarrow P_i$ : An instance of resource type  $R_j$  has been allocated to process  $P_i$ . (Assignment Edge).

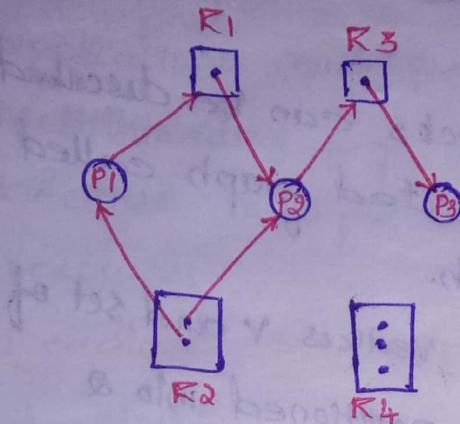
Pictorially:  $P_i$  is represented as circle &  $R_j$  as rectangle. Since  $R_j$  may have more than 1 instance, we represent each instance as a dot within rectangle. Note that the request edge points to only the rectangle  $R_j$ , whereas an assignment edge must also designate one of the dots in rectangle.

## Eg: RAG (Resource Allocation Graph)

$$P = \{P_1, P_2, P_3\}$$

$$R = \{R_1, R_2, R_3, R_4\}$$

$$E = \{P_1 \rightarrow R_1, P_2 \rightarrow R_3, R_1 \rightarrow P_2, R_2 \rightarrow P_2, R_2 \rightarrow P_1, R_3 \rightarrow P_3\}$$



Instances:

$$R_1 = 1$$

$$R_2 = 2$$

$$R_3 = 1$$

$$R_4 = 3$$

Process states:

✓ P1 is holding

Process	Holding	Waiting
P1	R2	R1
P2	R1, R2	R3
P3	R3	

Given a definition of the resource Allocation Graph, it can be shown that, if the graph contains no cycles then, no process in the system is deadlocked. If the graph does contain a cycle, then deadlock may exist.

Eg: Cycle → Deadlock: Consider the previous RAG, & suppose P3 requests an instance of resource type R2. Since no resource instance is currently available, a request edge  $P_3 \rightarrow R_2$  is added to the graph. At this point 2 minimal cycles exist in the system.

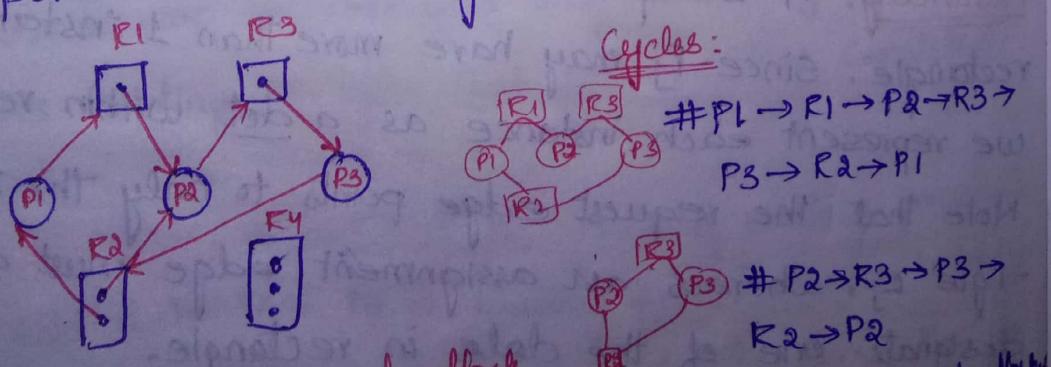


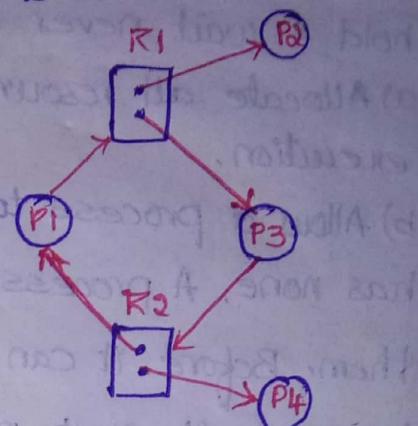
Fig: RAG with a deadlock

Processes P1, P2 & P3 are deadlocked

process P2 is waiting for R3, which is held by P3.  
P3 is waiting for either P1 or P2 to release R2. In  
addition, P1 is waiting for P3 to release R1.

eg: cycle → but no deadlock: There is no deadlock,  
because P4 process may release  
its instance of resource type R2.  
That resource can be allocated  
to P3, breaking the cycle, Then  
the system may or may not be  
still in a dead-locked state.

Cycle: P1 → R1 → P3 → R2 → P1



### Methods for Handling Deadlocks:

- 1) Deadlock Prevention: Provides a set of methods for ensuring that atleast 1 of the necessary conditions cannot hold.
- 2) Deadlock Avoidance: Requires that OS be given in advance additional inf concerning which resource a process will request & use during lifetime.
- 3) Deadlock Detection + Recovery: If above 2 methods are not used, then deadlock may occur. Then it shld be detected and recovered.

### ① Deadlock Prevention:

- 1) Mutual Exclusion: The mutual exclusion condition must hold for non-shareable resources. For example, a printer cannot be simultaneously shared by several processes.
- 2) Shareable resources, do not require ME access & thus cannot be involved in deadlock.
- 3) In general, we cannot prevent deadlocks by denying

the ME Condition, because some resources are intrinsically non-shareable.

ii) Hold + Wait: The following protocol can be used to ensure hold + wait never occurs.

a) Allocate all resources to a process before it begins execution.

b) Allow a process to request resources only when it has none. A process may request some resources & use them. Before it can request any additional resources, however, it must release all the resources that it is currently allocated.

#### Disadvantages:

1) Resource utilization may be low, since resources may be allocated but unused for a long period.

2) Starvation is possible: A process that needs several popular resources may have to wait indefinitely, because atleast 1 of the resources that it needs is always allocated to some other process.

3) No Preemption: The foll. protocol can be used: If a process is holding some resources & requests another resource that cannot be immediately allocated to it.

(that is the process must wait), then all resources the process is currently holding are preempted (released).

4) circular wait: One way to ensure that this condition never holds is to impose a total ordering of all resource types & to require that each process request in an increasing order of enumeration

Let  $R = \{R_1, R_2, R_3, \dots, R_m\}$  be the set of resource types. We assign a unique integer no. to each resource type, which allows us to compare 2 resources & to determine whether 1 precedes another in our ordering.

formally, we define a one-to-one function.

$F: R \rightarrow N$ , where  $N$  is set of Natural numbers.

e.g.:  $F(R_1) = 1$

$F(R_2) = 2$

$F(R_3) = 3$

$\vdots$

$F(R_m) = m$

Now, each process can request resources only in an increasing order of enumeration.

i.e., A process can request any no. of instances of resource types. Say  $R_i$ . After that, the process can request instances of resource type  $R_j$  if and only if  $F(R_j) > F(R_i)$ .

Alternatively, we can require that a process requesting an instance of resource  $R_j$  must have released any resources  $R_i$  such that  $F(R_i) \geq F(R_j)$ .

\* Deadlock Avoidance: Deadlock prevention may result in low device utilization & reduced sys. throughput. If the os has knowledge of what resources are requested by the processes before execution, then os can decide whether or not those resources are allocated to the processes to avoid possible deadlock in the future.

\* Each request requires that in making this decision the sys considers the resources currently available, the

resources currently allotted to each process, & the future requests & releases of each process.

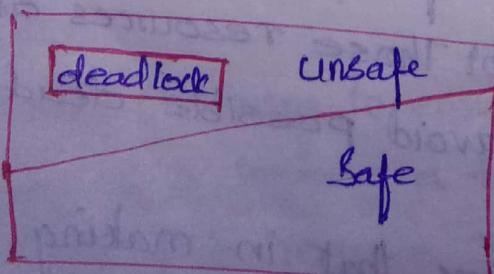
Safe State: A state is safe if the Sys can allocate resources to each process (upto its maximum) in some order & still avoid a deadlock.

More formally, a Sys is in a safe state only if there exists a safe sequence.

A sequence of processes  $\langle p_1, p_2, \dots, p_n \rangle$  is a safe sequence for the current allocation state if, for each  $p_i$ , the resource requests that  $p_i$  can still be satisfied by the currently available resources plus the resources held by all  $p_j$ , with  $j < i$ .

In this situation, if the resources that  $p_i$  needs are not immediately available, then  $p_i$  can wait until all  $p_j$  have finished. When they have finished,  $p_i$  can obtain all of its needed resources, complete its designated task, return its allocated resources, & terminate. When  $p_i$  terminates  $p_{i+1}$  can obtain its needed resources and so on. If no such sequence exists, then the Sys state is said to be unsafe.

\* A deadlocked state is an unsafe state & an unsafe state may lead to a deadlock.



consider 12 tape drivers + 3 processes.

<u>process</u>	<u>Max Need</u>	<u>Currently holding / allocated</u>
P0	10	5
P1	4	2
P2	9	2

(a) Find whether the sys. is in Safe state or not + also find the safe sequence?

$$\text{Total tape drivers} = 12$$

$$\text{Already allotted} = 5 + 2 + 2 = 9$$

$$\text{Free tape drivers} = 12 - 9 = 3$$

We need to use the available tape drives + find the safe sequence.

$$\text{safe sequence} = \langle P1, P0, P2 \rangle //$$

P1 needs 4 + it already has 2, allocate 2 more drives = 4, P1 can complete execution + it releases 4 drives. Now total drives available =  $4 + 1 = 5$ . Now we can allocate these 5 drives to P0  $\Rightarrow$  if release 10 drives, then P2 can be executed

chance of unsafe state: If P2 is allocated with 1 more

tape drive = ?

P1 can be executed but neither P0 nor P2 can complete their execution  $\Rightarrow$  Unsafe State.

Here idea is to keep the sys in a safe state.

Initially, the sys is in a safe state. Whenever a process requests a resource that is currently available, the sys must decide whether the resource can be allocated immediately or whether the process must wait. The request is granted only if the allocation leaves the sys in a safe state.

Resource Allocation Graph Alg: If we have a RA sys. with only 1 instance of each resource type, we can use a variant of the resource-allocation graph for deadlock avoidance.

✓ In addition to the request & assignment edges, a new type of edge called as claim edge is introduced. This is represented by a dashed line.

$P_i \dashrightarrow R_j \Rightarrow$  Process  $P_i$  may request resource  $R_j$  at sometime in the future.

✓ When the process  $P_i$  requests resource  $R_j$ , the claim edge  $P_i \dashrightarrow R_j$  is converted into a request edge  $P_i \rightarrow R_j$ . Similarly, when a resource  $R_j$  is released by  $P_i$ , the assignment edge  $R_j \rightarrow P_i$  is reconverted into a claim edge  $P_i \dashrightarrow R_j$ .

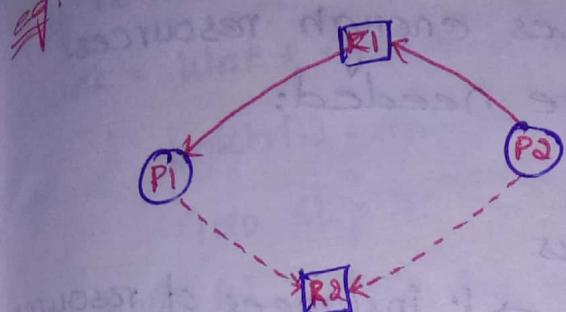
✓ The resources must be claimed by the process before starting the execution of process i.e., all claim edges must appear in the RAG.

✓ Now suppose that process  $P_i$  requests resource  $R_j$ . The resource can be granted only if converting the requesting edge  $P_i \rightarrow R_j$  to assignment edge  $R_j \rightarrow P_i$  does not result in the formation of cycle in RAG. We check for the safety by using a cycle-detection alg. Alg. may take an order of  $n^2$  operations (where  $n$  is the no. of processes in the Sys).

If no cycles exists, then the allocation of the resources will leave the sys in a safe state. If a cycle is found, then the allocation will put the

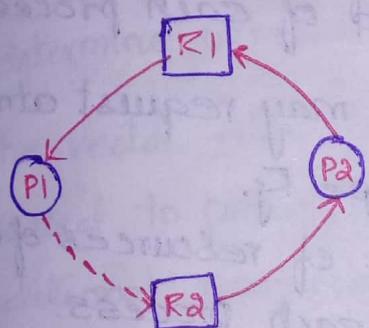
system in an unsafe state. In that case  $P_1$  will have to wait for its requests to be satisfied.

Consider the foll. situation.



will create a cycle in the graph.

Now suppose that  $P_2$  requests  $R_2$ . Although  $R_2$  is currently free, we cannot allocate it to  $P_2$ , since this allocation action



$P_1 \rightarrow R_2 \rightarrow P_2 \rightarrow R_1 \rightarrow P_1$

This indicates that the sys. is in an unsafe state. (If  $P_1$  requests  $R_2$  &  $P_2$  requests  $R_1$ , then a deadlock will occur).

Drawback: RAG alg. is not applicable to a resource allocation sys with multiple instances of each resource type. The next, Bankers algorithm is applicable to such sys. but less efficient than RAG.

Bankers Algorithm: The name was chosen because the alg. could be used in a banking sys to ensure that the bank never allocated its available cash in such away that it could no longer satisfy the needs of all its customers.

\* When a new process enters the sys, it must declare the max no. of instances of each resource type that it may need. Now, the sys will determine

whether the allocation of these resources will leave the sys in a safe state. If it will, the resources are allocated; otherwise the process must wait until some other process releases enough resources.

✓ the foll. datastructures are needed:

n: no. of processes

m: no. of resource types

1) Available:  $\text{Available}[j] = k \Rightarrow k$  instances of resource type  $R_j$  are available (Total resources)

2) Mark:  $n \times m$  matrix (max. demand of each process)

$\text{Mark}[i][j] = k \Rightarrow$  Process  $P_i$  may request atmost  $k$  instances of resource type  $R_j$ .

3) Allocation:  $n \times m$  matrix defines no. of resources of each type currently allocated to each process.

$\text{Allocation}[i][j] = k \Rightarrow$  Process  $P_i$  is currently allocated with  $k$  instances of resource type  $R_j$ .

4) Need:  $n \times m$  matrix defines the remaining resource need of each process.

$\text{Need}[i][j] = k \Rightarrow P_i$  may need  $k$  more instances of resource type  $R_j$  to complete its task.

Note:  $\text{Need}[i][j] = \text{Mark}[i][j] - \text{Allocation}[i][j]$ .

Safety Alg: The alg for finding out whether or not a sys is in a safe state.

Let work & finish be vectors of length  $m + n$  respectively.

Initialization: 1)  $\text{work} = \text{Available}$

$\text{Finish}[i] = \text{False}$  for  $i = 0, 1, 2, \dots, n-1$

2) Find an index  $i$  such that both

- a.  $\text{finish}[i] == \text{false}$ ;
- b.  $\text{Need}_i <= \text{Work}$
- if no such  $i$  exists goto step 4.
- 3)  $\text{Work} = \text{Work} + \text{Allocation};$
- $\text{finish}[i] = \text{true};$
- Goto Step 2.
- 4) if  $\text{finish}[i] == \text{true}$  for all  $i$ , then the Sys is in a safe state.

This alg. may require an order of  $m \times n^2$  operations to determine whether a state is safe.

- ✓ The vector Allocation; specifies the resources currently allocated to process  $P_i$ ; the vector Need $_i$  specifies the additional resources that process  $P_i$  may still request to complete its task.

\* Resource Request Alg: Alg to determine whether the requests can be safely granted or not.

- Let Request $_i$  be the request vector for process  $P_i$ . If  $\text{request}_{ij} = k$ , then process  $P_i$  wants  $k$  instances of resource type  $R_j$ . When a request for resources is made by process  $P_i$ , the following actions are taken.
- 1. If  $\text{Request}_i <= \text{Need}_i$ , goto Step 2. Otherwise, raise an error condition, since the process has exceeded its maximum claim.
- 2. If  $\text{Request}_i <= \text{Available}$ , goto Step 3. Otherwise,  $P_i$  must wait, since the resources are not available.
- 3. Have the system pretend to have allocated the requested

resources to process  $P_i$  by modifying the state as follows. (If 2 is true).

$$\text{Available} = \text{Available} - \text{Request}_i;$$

$$\text{Allocation} = \text{Allocation}_i + \text{Request}_i;$$

$$\text{Need} = \text{Need}_i - \text{Request}_i;$$

If the resulting resource allocation state is safe, the transaction is completed, & process  $P_i$  is allocated its resources. However if the new state is unsafe, then  $P_i$  must wait for  $\text{Request}_i$ , & the old resource-allocation state is restored.

Deadlock Detection: If a sys. does not employ either a deadlock-prevention or a deadlock-avoidance alg., then a deadlock situation may occur. In this environment the sys may provide:

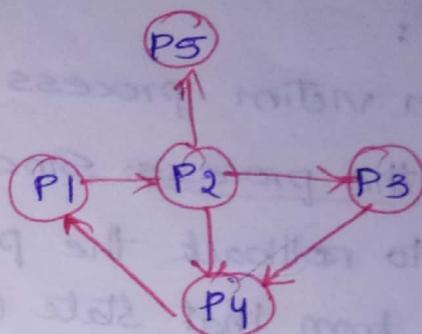
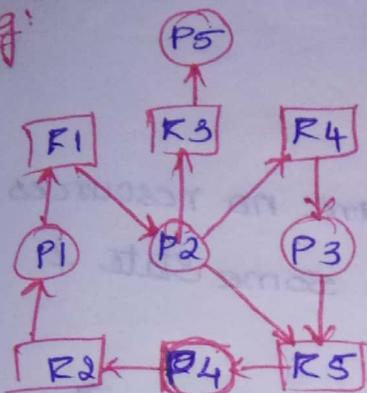
- 1) An alg. that examines the state of the sys to determine whether a deadlock has occurred.
- 2) An alg. to recover from the deadlock.

I. Single Instance of each Resource type: If all resources have only a single instance, then we can define a deadlock detection alg that uses a variant of the resource allocation graph, called a wait-for graph. We obtain this graph from RAG by removing the resource nodes & collapsing the appropriate edges.

More precisely, an edge from  $P_i$  to  $P_j$  in a wait-for graph implies that process  $P_i$  is waiting for process  $P_j$  to release a resource that  $P_i$  needs. As edge  $P_i \rightarrow P_j$  exists in a wait-for graph if & only if the corresponding

FAG graph contains 2 edges  $P_i \rightarrow R_q$  &  $R_q \rightarrow P_j$  for some resource  $R_q$ .  
✓ if there is a cycle  $\Rightarrow$  Deadlock.

e.g:



## II. Several instances of a Resource type.

✓ Bankers algorithm.

Recovery from Deadlock: When a deadlock detection alg determines that a deadlock exists, several alternatives are available. One possibility is to inform the operator that a deadlock has occurred & let the operator deal with the deadlock manually. Another possibility is to let the sys. recover from the deadlock automatically.

✓ There are 2 options for breaking a deadlock.

1) Process Termination

2) Resource Preemption.

1) Process Termination: 1) Abrt all deadlocked processes: will break the deadlock but at a great expense. The deadlocked processes may have computed for long time.

2) Abrt 1 process at a time until the deadlock cycle is eliminated: incurs a considerable overhead. After each process is aborted, a deadlock detection alg. must be invoked.

ii) Resource Preemption: Preempt some resources & give these resources to other processes until the deadlock cycle is broken.

If preemption is required, then 3 issues need to be addressed:

i) Selecting a victim process.

ii) Rollback the process: Since there are no resources we need to rollback the process to some state & restart it from that state (later).

iii) Starvation: If resources are preempted from the same process again & again then that process will be starved.